

Assignment 4

Introduction


Sebastian Dörrich (MSc)

Deep Learning Exercise

Chair of Explainable Machine Learning (xAI)



Overview




Overview




1. Convolutional Neural Networks

- a) CNN from scratch
- b) CNN with PyTorch

Convolutional Neural Networks



CNN from Scratch





ToDo

- Implement modules/layers:
 - Convolution
 - Max Pooling
 - ReLU
 - Linear
- Implement optimizer
 - SGD with momentum





Convolution (or cross-correlation)





Convolved feature

Image





Convolved feature

Image





Convolved feature

Image





Convolved feature

Image



Convolved feature

Image



Convolved feature

Image

<https://www.sciencedirect.com/topics/mathematics/convolutional-layer>

Vectorization

- It's okay to use loops for this assignment
- Vectorization is not necessary for this assignment but in practice, it is very important to speed up your code
- Note Convolutional layers are just linear transforms of kernels (weights) with different “views” of the image (i.e. toeplitz matrix)
- Use this notion to vectorize.

Vectorization – Tricks (1/2)

Goal:

Create views of an image the same size as the filters we're applying:

$$(c, h, w) \rightarrow (h_{\text{new}}, w_{\text{new}}, c, k, k)$$

Given an image A as a NumPy array:

$$A.\text{shape} = (c, h, w) = (3, 45, 40)$$

Assuming A is of type float64, its strides from NumPy are:

$$A.\text{strides} = (s_c, s_h, s_w) = (14400, 320, 8)$$

“Strides represent how many bytes of offset to get to the next value in that dimension.”

45×40×8 bytes per stride

40×8 bytes per stride

8 bytes per stride

Vectorization – Tricks (2/2)

In preparation for a convolution layer, we want to create “views” of the image with the same size as the filter. Assuming filter size $k=3$ and stride $\text{stride}=2$, we have:

$$h_{\text{new}} = (45 - 3) // 2 + 1 = 22$$

$$w_{\text{new}} = (40 - 3) // 2 + 1 = 19$$

```
np.lib.stride_tricks.as_strided(  
    A,  
    shape=(22, 19, 3, k, k),  
    strides=(s_h*2, s_w*2, s_c, s_h, s_w),  
    writeable=False)
```

Vectorization – Multiplying Tensors

As an example, given matrices:

A of shape $(n, c, h, w) \rightarrow$ containing n images

B of shape $(c, h, w, k) \rightarrow$ containing filters the same size as the image

How to apply the filters to all the images at once?

Method 1

`np.tensordot(A, B, axes=3)`

Method 2

`np.einsum('nchw,chwk->nk', A, B)`

Vectorization – Broadcasting

Given two matrices:

A is of size (a, b, c) and B is of size (b,)

adding them together along a specific dimension can be done via:

`A + B[:, np.newaxis]` or `A + B[:, None]`

This tells NumPy to broadcast across the 1st and 3rd dimensions before performing addition.

Vectorization – Unravel Index

How to retrieve the location of a point given the argmax?


```
A = np.array([[ 3,  5],  
              [17,  2],  
              [16,  3]])
```

```
i = A.argmax() = 2
```

This can be done via:

```
np.unravel_index(i, A.shape) = (1, 0)
```

CNN with PyTorch



ToDo

- Complete the training loop of PyTorch
 - Notice the difference between training and validation
- Implement two specified models
 - Two fully connected layers with a sigmoid activation function in between
 - Vanilla CNN




PyTorch – Complete the Training Loop

- Implement the training and validation step utilizing PyTorch
- Key concept is that in validation you won't be updating the gradients

PyTorch – Two Layer Model

Two fully connected layers with a sigmoid activation function in between



PyTorch – CNN

Implement a basic CNN utilizing PyTorch

